

## Importing Libraries And Datasets

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.tree import DecisionTreeClassifier, plot_tree
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score
```

```
df = pd.read_csv("/content/heart.csv")
df.head(10)
```

|   | age | sex | cp | trestbps | chol | fbs | restecg | thalach | exang | oldpeak | slope | ca |
|---|-----|-----|----|----------|------|-----|---------|---------|-------|---------|-------|----|
| 0 | 52  | 1   | 0  | 125      | 212  | 0   | 1       | 168     | 0     | 1.0     | 2     | 2  |
| 1 | 53  | 1   | 0  | 140      | 203  | 1   | 0       | 155     | 1     | 3.1     | 0     | 0  |
| 2 | 70  | 1   | 0  | 145      | 174  | 0   | 1       | 125     | 1     | 2.6     | 0     | 0  |
| 3 | 61  | 1   | 0  | 148      | 203  | 0   | 1       | 161     | 0     | 0.0     | 2     | 1  |
| 4 | 62  | 0   | 0  | 138      | 294  | 1   | 1       | 106     | 0     | 1.9     | 1     | 3  |
| 5 | 58  | 0   | 0  | 100      | 248  | 0   | 0       | 122     | 0     | 1.0     | 1     | 0  |
| 6 | 58  | 1   | 0  | 114      | 318  | 0   | 2       | 140     | 0     | 4.4     | 0     | 3  |
| 7 | 55  | 1   | 0  | 160      | 289  | 0   | 0       | 145     | 1     | 0.8     | 1     | 1  |
| 8 | 46  | 1   | 0  | 120      | 249  | 0   | 0       | 144     | 0     | 0.8     | 2     | 0  |
| 9 | 54  | 1   | 0  | 122      | 286  | 0   | 0       | 116     | 1     | 3.2     | 1     | 2  |

Next steps:

[Generate code with df](#)[New interactive sheet](#)

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1025 entries, 0 to 1024
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype  
---  -
 0   age         1025 non-null  int64  
 1   sex         1025 non-null  int64  
 2   cp          1025 non-null  int64  
 3   trestbps    1025 non-null  int64  
 4   chol        1025 non-null  int64  
 5   fbs         1025 non-null  int64  
 6   restecg     1025 non-null  int64  
 7   thalach     1025 non-null  int64
```

```
8  exang    1025 non-null  int64
9  oldpeak  1025 non-null  float64
10 slope    1025 non-null  int64
11 ca       1025 non-null  int64
12 thal     1025 non-null  int64
13 target   1025 non-null  int64
dtypes: float64(1), int64(13)
memory usage: 112.2 KB
```

```
df.describe()
```

|              | age         | sex         | cp          | trestbps    | chol        | fbs         |
|--------------|-------------|-------------|-------------|-------------|-------------|-------------|
| <b>count</b> | 1025.000000 | 1025.000000 | 1025.000000 | 1025.000000 | 1025.000000 | 1025.000000 |
| <b>mean</b>  | 54.434146   | 0.695610    | 0.942439    | 131.611707  | 246.000000  | 0.149268    |
| <b>std</b>   | 9.072290    | 0.460373    | 1.029641    | 17.516718   | 51.59251    | 0.356527    |
| <b>min</b>   | 29.000000   | 0.000000    | 0.000000    | 94.000000   | 126.000000  | 0.000000    |
| <b>25%</b>   | 48.000000   | 0.000000    | 0.000000    | 120.000000  | 211.000000  | 0.000000    |
| <b>50%</b>   | 56.000000   | 1.000000    | 1.000000    | 130.000000  | 240.000000  | 0.000000    |
| <b>75%</b>   | 61.000000   | 1.000000    | 2.000000    | 140.000000  | 275.000000  | 0.000000    |
| <b>max</b>   | 77.000000   | 1.000000    | 3.000000    | 200.000000  | 564.000000  | 1.000000    |

```
df.shape
```

```
(1025, 14)
```

```
df.isnull().sum()
```

|                 | 0 |
|-----------------|---|
| <b>age</b>      | 0 |
| <b>sex</b>      | 0 |
| <b>cp</b>       | 0 |
| <b>trestbps</b> | 0 |
| <b>chol</b>     | 0 |
| <b>fbs</b>      | 0 |
| <b>restecg</b>  | 0 |
| <b>thalach</b>  | 0 |
| <b>exang</b>    | 0 |
| <b>oldpeak</b>  | 0 |
| <b>slope</b>    | 0 |
| <b>ca</b>       | 0 |
| <b>thal</b>     | 0 |
| <b>target</b>   | 0 |

**dtype:** int64

## ✓ Training A Decision Tree Classifier

```
X = df.drop('target', axis=1)
y = df['target']
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_
```

```
print(f"Training Set Size: {X_train.shape[0]}")
print(f"Testing Set Size: {X_test.shape[0]}")
```

```
Training Set Size: 717
Testing Set Size: 308
```

```
dt = DecisionTreeClassifier(random_state=42)
dt.fit(X_train, y_train)

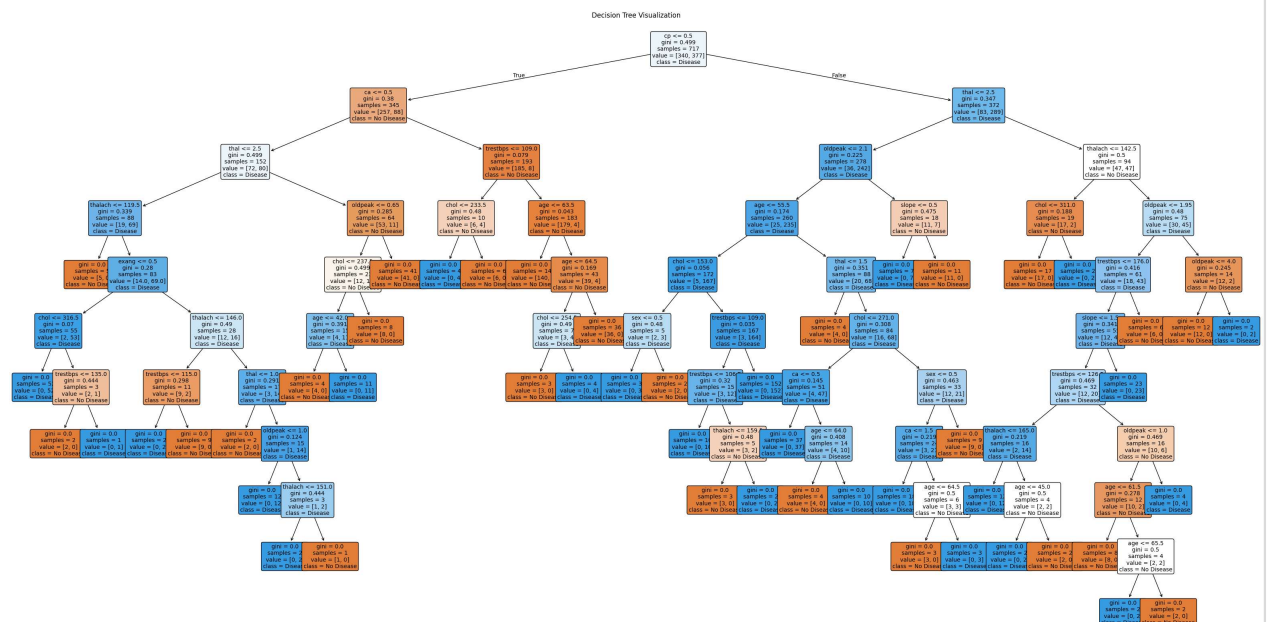
y_train_pred_dt = dt.predict(X_train)
y_test_pred_dt = dt.predict(X_test)

train_acc_dt = accuracy_score(y_train, y_train_pred_dt)
test_acc_dt = accuracy_score(y_test, y_test_pred_dt)
```

```
print(f"Decision Tree (Unconstrained) Training Accuracy: {train_acc_dt:.5f}")
print(f"Decision Tree (Unconstrained) Testing Accuracy: {test_acc_dt:.5f}")
```

Decision Tree (Unconstrained) Training Accuracy: 1.00000  
Decision Tree (Unconstrained) Testing Accuracy: 0.97078

```
plt.figure(figsize=(40, 20))
plot_tree(
    dt,
    feature_names=X.columns.tolist(),
    class_names=['No Disease', 'Disease'],
    filled=True,
    rounded=True,
    fontsize=10
)
plt.title("Decision Tree Visualization")
plt.show()
```



## ✧ Controlling Tree Depth

```
dt_pruned = DecisionTreeClassifier(max_depth=4, min_samples_split=10, random_state=42)
dt_pruned.fit(X_train, y_train)

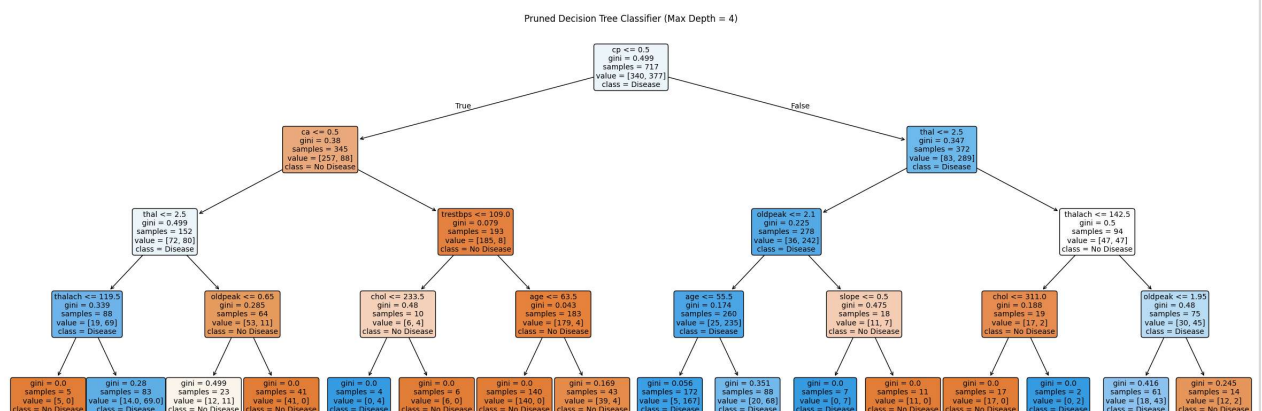
y_train_pred_pruned = dt_pruned.predict(X_train)
y_test_pred_pruned = dt_pruned.predict(X_test)

train_acc_pruned = accuracy_score(y_train, y_train_pred_pruned)
test_acc_pruned = accuracy_score(y_test, y_test_pred_pruned)

print(f"Decision Tree (Pruned) Training Accuracy: {train_acc_pruned:.5f}")
print(f"Decision Tree (Pruned) Testing Accuracy: {test_acc_pruned:.5f}")
```

Decision Tree (Pruned) Training Accuracy: 0.89679  
Decision Tree (Pruned) Testing Accuracy: 0.83442

```
plt.figure(figsize=(30, 10))
plot_tree(
    dt_pruned,
    feature_names=X.columns.tolist(),
    class_names=['No Disease', 'Disease'],
    filled=True,
    rounded=True,
    fontsize=10
)
plt.title("Pruned Decision Tree Classifier (Max Depth = 4)")
plt.show()
```



## ✧ Training A Random Forest

```
rf_classifier = RandomForestClassifier(n_estimators=100, random_state=42, max_depth=4)
rf_classifier.fit(X_train, y_train)
```

```
y_train_pred_rf = rf_classifier.predict(X_train)
y_test_pred_rf = rf_classifier.predict(X_test)

train_acc_rf = accuracy_score(y_train, y_train_pred_rf)
test_acc_rf = accuracy_score(y_test, y_test_pred_rf)

print(f"Random Forest Training Accuracy: {train_acc_rf:.5f}")
print(f"Random Forest Testing Accuracy: {test_acc_rf:.5f}")
```

Random Forest Training Accuracy: 0.93305

Random Forest Testing Accuracy: 0.88636

## ✓ Interpreting Feature Importances

```
feature_importances = pd.Series(rf_classifier.feature_importances_, index=X.columns)
sorted_features = feature_importances.sort_values(ascending=False)

print("Feature Importances In Random Forest:")
print(sorted_features.head(20))

plt.figure(figsize=(10, 6))
sorted_features.plot(kind='bar')
plt.title("Random Forest Feature Importances")
plt.ylabel("Importance Score")
plt.show()
```


Feature Importances In Random Forest:

```

cp          0.179715
thal        0.168812
oldpeak     0.137887
ca          0.134040
thalach     0.112034
exang       0.060258
slope       0.057825
age         0.052019
trestbps    0.033680
chol        0.027935
sex         0.025397
restecg     0.006950
fbs         0.003448
dtype: float64

```

Random Forest Feature Importances



## ▼ Cross Validation

```

cv_scores_dt = cross_val_score(dt_pruned, X, y, cv=5, scoring='accuracy')

cv_scores_rf = cross_val_score(rf_classifier, X, y, cv=5, scoring='accuracy')

print(f"Decision Tree CV Scores: {cv_scores_dt}")
print(f"Decision Tree Mean CV Accuracy: {cv_scores_dt.mean():.5f}")
print(f"Random Forest CV Scores: {cv_scores_rf}")
print(f"Random Forest Mean CV Accuracy: {cv_scores_rf.mean():.5f}")

```

```

Decision Tree CV Scores: [0.84390244 0.83414634 0.87317073 0.80487805 0.81463415]
Decision Tree Mean CV Accuracy: 0.83415

```